

**PARQ: A MEMORY-EFFICIENT APPROACH FOR  
QUERY-LEVEL PARALLELISM**

A Thesis Presented

by

QIANQIAN GAO

Submitted to the Graduate School of the  
University of Massachusetts Amherst in partial fulfillment  
of the requirements for the degree of

MASTER OF SCIENCE IN ELECTRICAL AND COMPUTER ENGINEERING

July 2016

Department of Electrical and Computer Engineering

© Copyright by Qianqian Gao 2016  
All Rights Reserved

**PARQ: A MEMORY-EFFICIENT APPROACH TO REALIZE  
QUERY-LEVEL PARALLELISM**

A Proposal Presented

by

QIANQIAN GAO

Approved as to style and content by:

---

Lixin Gao, Chair

---

Michael Zink, Member

---

David Irwin, Member

---

C. V. Hollot, Department Chair  
Department of Electrical and Computer  
Engineering

# **ABSTRACT**

## **PARQ: A MEMORY-EFFICIENT APPROACH TO REALIZE QUERY-LEVEL PARALLELISM**

September 2016

QIANQIAN GAO

B.S., BEIJING UNIVERSITY OF POSTS AND TELECOMMUNICATIONS

M.S., UNIVERSITY OF MASSACHUSETTS AMHERST

Directed by: Professor Lixin Gao

In the era of big data, people not only enjoy what massive information brings, but also experience the problem of information overload. As the volumes of both data and users are increasing sharply, more and more studies focus on how to answer a query for interesting information from massive data. However, most query systems are designed and implemented to optimize the performance in processing a single query and do not support data sharing among query processing jobs. When they are extended to process multiple concurrent queries, they will suffer the problems of the inefficient use of memory and waste of time.

This proposal aims to design and implement a memory-efficient system, ParQ, which could be adopted by query systems to realize query-level parallelism. The main idea includes constructing a common memory block for maintaining sharable data. By sharing data, ParQ is able to process multiple queries concurrently while reducing memory usage and running time. We applied several to several existing query systems. The experiment results show that ParQ improves the performance in both job completion time and memory usage, when executing multiple concurrent query jobs.

## TABLE OF CONTENTS

|  |           |
|--|-----------|
| <b>ABSTRACT</b> .....  | <b>4</b>  |
| <b>CHAPTER 1</b> .....   | <b>7</b>  |
| <b>INTRODUCTION</b> .....  | <b>7</b>  |
| 1.1    Trends and challenges.....  | 7         |
| 1.2    Motivation.....   | 8         |
| 1.3    Contribution .....  | 8         |
| 1.4    Proposal outline.....   | 9         |
| <b>CHAPTER 2</b> .....   | <b>10</b> |
| <b>PROBLEM DEFINITION</b> .....  | <b>10</b> |
| <b>CHAPTER 3</b> .....   | <b>13</b> |
| <b>DESIGN AND IMPLEMENTATION OF PARQ</b> .....                               | <b>13</b> |
| 3.1    Overall design .....  | 13        |
| 3.2    System implementation .....   | 16        |
| 3.2.1    API for data loading .....  | 16        |
| 3.2.2    API for query processing .....                                      | 17        |
| 3.2.3    Ability to adapt to distributed environment.....                    | 19        |
| 3.3    Instruction for application.....                                      | 21        |
| <b>CHAPTER 4</b> .....   | <b>22</b> |
| <b>APPLICATION ON EXISTING SYSTEMS</b> .....                                 | <b>22</b> |
| 4.1    Application on The top-k Personalized PageRank system .....           | 22        |
| 4.1.1    System specification .....  | 22        |
| 4.1.2    Implementation process .....  | 23        |
| 4.2    Application on the top-k star query system.....                       | 25        |
| 4.2.1    System specification .....  | 25        |
| 4.2.2    Implementation process .....  | 26        |
| 4.3    Application on the belief propagation system .....                    | 28        |
| 4.3.1    System specification .....  | 28        |
| 4.3.2    Implementation process .....  | 29        |
| <b>CHAPTER 5</b> .....   | <b>32</b> |
| <b>APPLICATION AND EVALUATION</b> .....                                      | <b>32</b> |
| 5.1    Overview .....  | 32        |
| 5.2    Performance of the top-k Personalized PageRank system using ParQ..... | 32        |

|                                   |   |           |
|-----------------------------------|---|-----------|
| 5.2.1                             | Experiment settings .....                                     | 32        |
| 5.2.2                             | Running time .....  | 33        |
| 5.2.3                             | Memory usage .....  | 34        |
| 5.3                               | Performance of the top- $k$ star query system using ParQ..... | 34        |
| 5.3.1                             | Experiment settings .....                                     | 35        |
| 5.3.2                             | Running time .....  | 35        |
| 5.3.3                             | Memory usage .....  | 36        |
| 5.4                               | Performance of the belief propagation system.....             | 37        |
| 5.4.1                             | Experiment settings .....                                     | 37        |
| 5.4.2                             | Running time .....  | 37        |
| 5.4.3                             | Memory usage .....  | 38        |
| <b>CHAPTER 6</b>                  | .....   | <b>40</b> |
| <b>CONCLUSION AND FUTURE WORK</b> | .....   | <b>40</b> |
| <b>CHAPTER 7</b>                  | .....   | <b>41</b> |
| <b>RELATED WORK</b>               | .....   | <b>41</b> |
| 7.1                               | Large structured dataset processing.....                      | 41        |
| 7.2                               | Algorithm implementations .....                               | 41        |
| <b>References</b>                 | .....   | <b>42</b> |

## CHAPTER 1

### INTRODUCTION

#### 1.1 Trends and challenges

The future of big data is promising, even though many challenges remain. Data has become a class of asset. In the areas of Internet, Telecommunications, Finance, and etc., many companies have achieved commercial success by using big data.

On the other hand, the development of big data also faces many challenges, for example, how to reduce data fragment. Nowadays, data is organized into different databases that may belong to different organizations. Even in the same company, data may be maintained by different departments. It is necessary to integrate all the data to better exploit the potential of information.

How to improve data preprocessing is also a big challenge. At the step of data preprocessing, data is cleaned, denoised and transformed into an easy-to-use structure. Nonstandard preprocessing often leads to inaccuracies, poor quality and low availability. For big data area, not only data collecting is important but also preprocessing.

Another challenge is how to find requested information from massive data, especially as the volumes of both data and users are increasing sharply. According to recent statistics, by August 2014, Google Knowledge Graph had contained 1.6 billion labeled entities [1]. As of January 2014, Freebase had approximately 44 million topics and 2.4 billion facts [2]. To make use of such databases, many researchers focus on algorithms of data processing. Topics include, for example, how to improve accuracy, how to increase efficiency, and how to reduce cost, such as memory usage.

## 1.2 Motivation

Query systems refers to an implementation that aims to answer queries on a database. Such applications have been widely applied in the areas of search engine, social network, e-commerce, and etc. In practical applications, query systems may be queried by more than one user at the same time so that they are supposed to process multiple queries concurrently. The job execution logs from a large Chinese social network shows that at peak time, there are more than 20 jobs submitted to the platform [3]. Some developing query systems are inherently designed to evaluate the the performance in answering a single query and difficult to expended to process multiple queries. Therefore, although such query systems are efficient when processing a single query, they require huge memory usage and long processing time when they process a number of queries concurrently. This is because they cannot support multiple jobs to share in-memory data and each job has to keep an exclusive copy of the same graph data, which results in a serious waste of time and memory.

This proposal aims to seek an approach to enable such systems to answer multiple queries concurrently by sharing sharable data among different queries. This approach should ensure 1) a query system uses only one copy of sharable data when it processes multiple queries, 2) the processing time of a query will not be affected by other queries, and 3) this approach should be easily used by different query systems.

## 1.3 Contribution

This proposal introduces a memory-efficient approach, ParQ system, to enable query systems to answer multiple queries concurrently by sharing data among different query jobs. In this proposal, the sharable data is the dataset from which a query system finds answers of a query.

The ParQ system offers a standard interface. It provides the ability to automatically process data in different formats as long as the formats conform to the system's requirements. This allows ParQ to be easily used by many query systems with a little manual work.

Besides the applications on the query systems, ParQ system could also apply to other similar situations where different processes require the same data.

As case studies, we apply the ParQ system to two query systems. A series of experiments are performed on a real computer cluster, and ParQ system shows good performance. It achieves approximately  $n$  times less memory usage to answer  $n$  queries concurrently than existing query systems. When there is limited memory, ParQ also enables more concurrent jobs and less processing time.

#### **1.4 Proposal outline**

In Chapter 2, we give an explicit problem definition. Chapter 3 describes the overall design of the proposed approach, gives a detailed explanation of each method provided, and briefly introduce how to use the ParQ system. In Chapter 4, we apply ParQ to one real query systems. In Chapter 5, we design and perform a series experiments to evaluate the performance when using ParQ. Chapter 6 introduces the future work in the following month and Chapter 7 gives the introduction to related work.

## CHAPTER 2

### PROBLEM DEFINITION

With the great increase of data volume in information area, more and more applications focus on mining interesting information from massive data. Some classes of massive data are normally stored into structured datasets to be further used. For example, Freebase is an online collection of structured data. It is a global resource, which allows people to access information. A query is a request for information on a target dataset. A query on Freebase can be “When is the date of Kennedy’s birthday?” Query systems are implementations that adopt a specific data processing algorithm and aim to answer queries on a target dataset. For a query system, a job refers to a process that is responsible for answering a single query. In the process of answering queries on the same dataset, the data used can be divided into two parts: static data and dynamic data. Static data is used but not changed. In this proposal, static data refers to the target dataset, because almost jobs are executed on the same target dataset. For example, Facebook performs most computations on the same friendship graph [Seraph 28]. Dynamic data refers to the data generated during the computation. It is query-specific and constantly updated. For example, the rank score of a candidate answer is the dynamic data. Therefore a query answering process could be viewed as a computation using a set of static data and a set of dynamic data, as shown in Figure 1.

Traditionally, large data is stored in disk, such as database. However, in order to read data faster, a number of query systems load static data into memory. Since many query systems are inherently designed for evaluating the performance of answering a single query, they usually tightly combine static data and data together. Therefore, they do not allow multiple jobs to share the static data in memory and each individual job has to maintain an exclusive copy of the static data in memory even if the static data is the same. Such query systems are much inefficient when they try to process

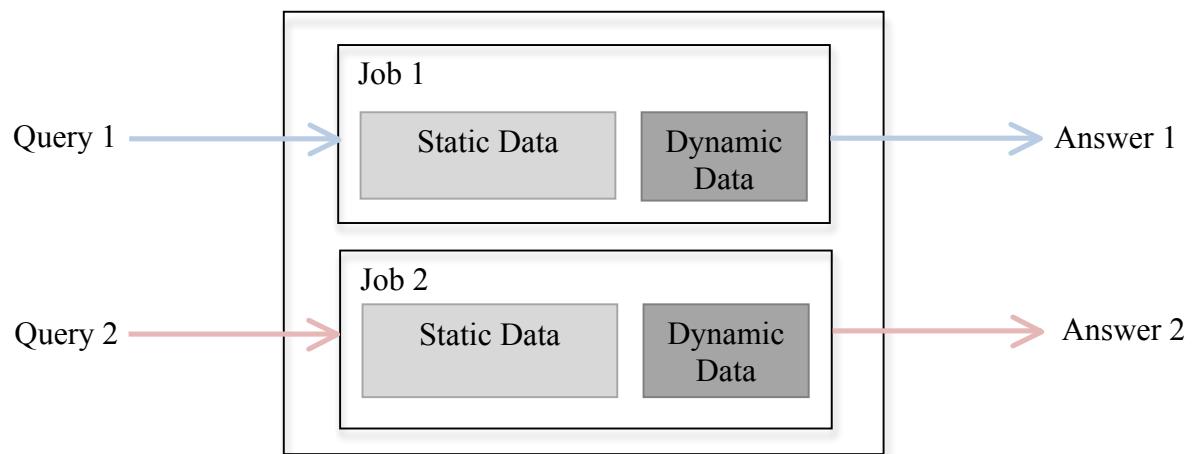
multiple queries. The inefficiencies mainly fall in two aspects: longer job completion time and large memory usage.

To be more specific, there are two approaches to answer  $n$  concurrent queries. First, a query system could keep one copy of the static data in memory and process the  $n$  queries in serial. In this situation, if the running time for completing one query is  $t$ , the total running time is  $t*n$ . Second, the system could also keep  $n$  copies of the static data in memory and start  $n$  jobs to answer  $n$  queries in parallel, as shown in Figure 2. In this case, if the memory usage for completing one query is  $m$ , the total memory usage for completing  $n$  queries is  $m*n$ , which could be a huge number. When there is limited memory, the number of running jobs is also limited and the large memory usage also results in longer running time.

The objective in this proposal is to seek a memory-efficient approach to enable a query system to answer multiple queries concurrently. One possible way is to share static data among different jobs. This approach should also ensure that 1) a query system keeps only one copy of static data among different jobs, 2) in the ideal situation, the processing time of one job will not be affected by other jobs, and that 3) it can be adopted by different query systems where datasets, queries, and implementations are all different.



**Figure 1: A job of a query system**



**Figure 2: Answering concurrent queries with existing systems**

## CHAPTER 3

### DESIGN AND IMPLEMENTATION OF PARQ

#### 3.1 Overall design

The objective in this proposal is to seek a memory-efficient approach to enable a query system to concurrently answer multiple queries. Since almost all concurrent jobs use the same static data, the main idea of the proposed approach, the ParQ system, is to create a shared memory to realize static data sharing among concurrent query processing jobs.

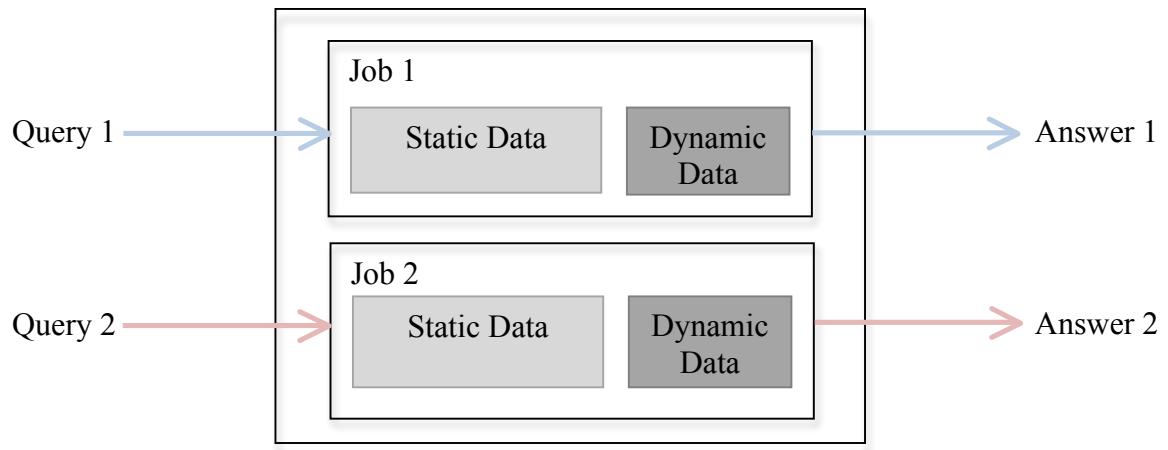
As explained in Chapter 2, a query system can start  $n$  jobs to answer  $n$  queries in parallel with  $n$  copies of the static data in memory. As shown in the first box in Figure 3, when a query system processes two queries on the same static dataset at the same time, two jobs are started and each job is responsible for one query. Each job accepts one query, loads a copy of static data, performs computations, returns answers, and terminates. As a result, the query system has to keep two copies of the static data in memory.

In order to share static data, the design of ParQ splits the task in an original job into two parts: data loader and query processor. A job of data loader creates a shared memory and stores the static data into the memory. A job of query processor accesses the shared memory and fetches requested data. ParQ starts  $n$  jobs of query processor to process  $n$  queries. Consider as Figure 3 (b) as the example. When processing two queries on the same static dataset, ParQ first starts a job of data loader, Job 0, to create a shared memory and store the static data in the shared memory. Then, it starts two jobs of query processor, Job 1 and Job 2, to handle the two queries. Job 1 and Job 2 can connect to the shared memory to get needed data. In this way, ParQ only needs to keep one copy of the static data in memory.

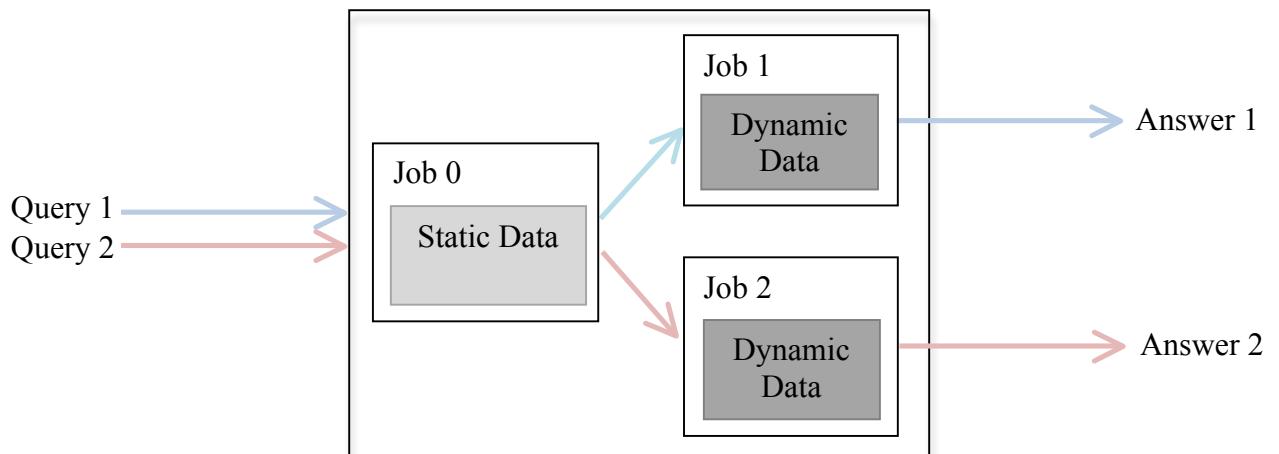
There are two main challenges in implementing ParQ. First of all, it should be easy to apply ParQ to different query systems where datasets, queries, and implementations are all different.

Different datasets have different data structures. ParQ should support diverse static data structures in order to be applied to different query systems with little manual work.

Secondly, ParQ should be able to work in a distributed environment. The shared memory storing the static data may be maintained by more than one worker. Therefore, the segments of the shared memory should be indexed properly to avoid potential errors.



(a) Query system keeps n copies of the static data



(b) ParQ keeps one copy of the static data

Figure 3: Multiple queries processing

The ParQ system contains two classes, one for data loader to process the static data and the other one for query processor to fetch requested data. It can be applied to different query systems and be able to work in a distributed environment. The following sections will introduce the api provided by ParQ, including the implementation and usage.

## 3.2 System implementation

As described in the last section, a data loader of the ParQ system enables static data sharing among different jobs. Each query processor of ParQ is responsible for one job and can connect to the shared memory to get needed data.

This section introduces the implementation of data loader and query processor and explains how to use ParQ system.

### 3.2.1 API for data loading

Data loader creates a shared memory and stores the target static data into the memory. Data loader stores every number in the specific data type and assigns one label to each number. It can support different data formats as long as the formats conform to several requirements.

The ParQ system provides a class, `DataLoader`, as shown in Figure 4. Given the input file path, the function `loadData(string path)` processes the data in the input file.

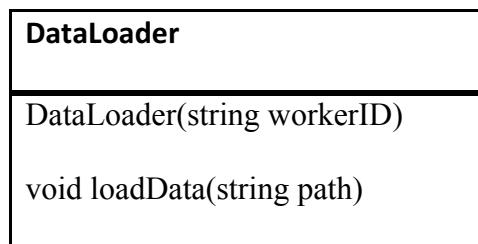


Figure 4: `DataLoader` class

In order to atomically process different files, ParQ requires the input files to be in a specific format. First, as Line c and Line d in Figure 5, a line represents a set of data and each tap-separated column contains a data item. The data in different lines should be organized in the same sequence, namely, the data in the same column has the same meaning and data type. The first column must be the index. An index is an integer and can specify one line. The following columns can be any data whose data type is supported. At present, the ParQ system supports “int”, “float”, “string”,

“vector<int>”, “vector<float>”, and “vector<string>”. Among them, the numbers in vectors are separated by spaces.

Secondly, users should assign labels to the columns in the first line, as Line a in Figure 5, and specify the data types in the second line as Line b in Figure 5. Since the data items in the first column are integer indices, users don’t need to specify the first column.

For an input file that meets all the requirements, the data loader of ParQ will store all data into the shared memory in the specific data type and name the data by its index and label. For example, for an input file demonstrated in Figure 6, ParQ will process “-32959029” as an integer, store it into the shared memory, and name it as “45-NodeType”.

|   |        |        |        |
|---|--------|--------|--------|
| a |        | label1 | label2 |
| b |        | type1  | type2  |
| c | index1 | data1  | data2  |
| d | index2 | data3  | data3  |

**Figure 5: Input file requirements**

|     | NodeType  | Neighbor                    |
|-----|-----------|-----------------------------|
|     | int       | vector<int>                 |
| 45  | -32959029 | 114661706 131784151         |
| 476 | -32959029 | 532524246 68210119 20109426 |

**Figure 6: An example of input file**

### 3.2.2 API for query processing

After loading the static data, users can start query processors to process query. The ParQ system provides the QueryProcessors class, as shown in Figure 7, for query processors to access the shared memory and fetch needed data.

**QueryProcessor**

```
QueryProcessor(string workerID)

int getInt(string index, string label)

float getFloat(string index, string label)

string getString(string index, string label)

int getFromVectorInt(string index, string label, string vectorIndex)

float getFromVectorFloat(string index, string label, string vectorIndex)

string getFromVectorString(string index, string label, string vectorIndex)
```

**Figure 7: QueryProcessor class**

To fetch the target data, the users need to provide the data type, the index, and the label.

Consider the input file shown in Figure 8 as the example. If a user tries to fetch the node type of node 45 from the shared memory, he/she should provide the data type and the label of the target. Since the data type is integer, the label is “NodeType”, and the index is 45, the user could call the function getInt(“45”, “NodeType”). This function will connect to the shared memory, search for the data item, and return “-32959029”. Similarly, the user can get a float, or a string from the shared memory by using corresponding function.

If the target data is in a vector, users should also provides the vector index. For example, a user wants to get the first neighbor node of node 45. He/she should call function getFromVectorInt(string index, string label, string vectorIndex), namely getFromVectorInt(“45”, “Neighbor”, “1”). The user could get an item for a vector of floats, a vector of doubles or a vector of strings in the same way.

|     | NodeType  | Neighbor                    |
|-----|-----------|-----------------------------|
|     | int       | vector<int>                 |
| 45  | -32959029 | 114661706 131784151         |
| 476 | -32959029 | 532524246 68210119 20109426 |

Figure 8: An example of input file

### 3.2.3 Ability to adapt to distributed environment

The ParQ system starts a data loading job and  $n$  query processing jobs to process  $n$  queries. In distributed model, a job initializes a master and a certain number of works. The master is responsible for work scheduling and the workers share the tasks of the job.

Consider Figure 9 as the example. There are also two workers for each job. Job 0 is a data loading job and aims to load static data. Each worker of Job 0 maintains a part of the static data. Job 1 and Job 2 are query processing jobs and aim to process queries. Each worker of them shares the computations and keeps a portion of dynamic data. The dynamic data is generated during computations.

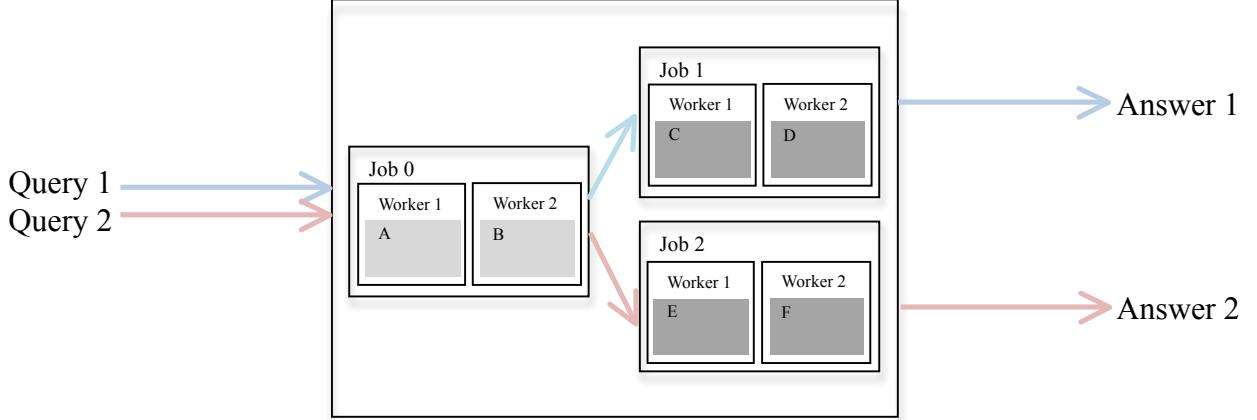


Figure 9: Two jobs in distributed model when ParQ applied

In distributed model, a large static dataset is split into several partitions based on partitioning strategies. Work  $i$ s are responsible for Part  $i$ . To be more specifically, Worker  $i$  of data loader creates a shared memory on a physical machine and loads Part  $i$  into the memory. Worker  $i$ s of query processors must be on the same physical machine to access Part  $i$ . As shown in Figure 9, A, C, and E must be on the same physical machine, so that C and E can access the shared memory creates by A. ParQ ensures Worker  $i$ s are assigned to the same machine by creating a map between worker ID and machine ID.

It is possible that workers with different ID are assigned to the same physical machine. In this case, by giving identifiers to data loader and query processors based on worker ID, ParQ ensures that Worker  $i$ s of query processors connects to the shared memory created by Worker  $i$  of data loader. Consider Figure 9 as the example, A, B, C, D, E, and F are assigned to the same physical machine. The memory created by A is named by “1” and the memory created by B is named by “2”. The query processor tasks on C and E are given an identifier “1”, so that C and E knows they should connects to the shared memory “1”.

### 3.3 Instruction for application

The ParQ system can be easily used by many query systems with a little manual work. Figure 10 describes the pseudo-code to apply ParQ. The static data of a query system should be presented in required format. The task in an original job should be split into two parts: data loading and query processing. A job uses an *if* statement to determine whether itself is a data loader or a query processor. If it is a data loader, it will create a shared memory and store the static data into the shared memory. If it is a query processor, it will access the shared memory, fetch requested data, and perform computations. The following example applications show how ParQ can be applied to a real query system.

*Input:* static data in required format, queries

```
if(current job is data loader) {  
    create an instance of DataLoader;  
    call corresponding functions  
    terminate  
} else if(current job is query processor) {  
    create an instance of QueryProcessor;  
    call corresponding functions  
    terminate  
}
```

**Figure 10: The high-level architecture to apply ParQ**

## CHAPTER 4

### APPLICATION ON EXISTING SYSTEMS

#### 4.1 Application on The top-k Personalized PageRank system

The ParQ system is first applied to the top-k PPR system. In the following sections, after a brief introduction of the top-k PPR system, we will show how to apply ParQ to it.

##### 4.1.1 System specification

Personalized recommendation is a key task for various applications. The top- $k$  PPR system makes personalized recommendation by implementing the Personalized PageRank computation. It aims to find  $k$  items that are the most relevant to a query set from an item graph. The query item set is formed based on known personal preferences. Personalized PageRank is used to measure the relevance of candidates by computing the stationary distribution of a random walk. In each computing step, the random walk reaches a random out-neighbor with probability  $d$  or jumps to a query node with probability  $1 - d$ .

The top- $k$  PPR system is an extension of the maiter framework [5]. In this system, relevance score is computed by accumulative computation in an asynchronous manner. A fraction of nodes are selected to update their scores in each pass. The chosen nodes will update its value ( $v$ ) and propagate the change of value ( $\Delta v$ ) to their neighbors. The top- $k$  PPR system uses score bounds to speed up the computation.

To process a query, the top- $k$  PPR system will start a process, or job, that loads the graph into memory and performs an accumulative computation on the in-memory graph to measure the relevance of candidates. The computation is performed until the top- $k$  answers are obtained.

#### 4.1.2 Implementation process

As described in the last section, the top- $k$  PPR system starts a process, or job, to process a query. The job loads the graph into memory and performs an accumulative computation on the in-memory graph to measure the relevance of candidates. The computation is performed until the top- $k$  answers are obtained.

When answering multiple queries on the same graph, the item graph used is static and can be shared among different jobs. Therefore, we apply the ParQ system to enable in-memory data sharing among query processing jobs.

First, Figure 11 shows how an item graph should be represented in a text file to use ParQ. Each line describes one node. The first column is an integer representing the node name or the node key. The next two are the nodes attributes used in bound computation. The following column lists the node's neighbors. As explained in Section 3.2.1, the data types of the columns except the first should be specified in the first two lines.

```

globalScore    incomingDegree  Neighbors
float          float          vector<int>
0  30.349953    138.460022    11342 824020 867923 891835
1  0.38410074    7.46896818    53051 203402 223236 276233 552600 569212 635575 748615 862566 893884
2  3.7941647    15.35852867    30957 357310 423174 430119 462435 472889 565424 581609 597621 644135 858904
3  0.8809524    2.466473115   87562 131116 262285 290424 298887 449136 456686 718023 789669 812470

```

**Figure 11: A part of an item graph for the top- $k$  PPR system**

Second, we write the code of data loader. In the original system, method `read_graph` in *struct PPRGlobalIterateKernel* is responsible for data loading. To apply ParQ, we create an instance of *DataLoader* and pass in the current worker ID as shown in Figure 12. Then, we call the method `loadData` to enable in-memory static data sharing.

```

void read_graph(string& partition_file, TypedGlobalTable<int, double, double, PPRGlobal_data *> *table) {
    ...
    DataLoader loader(workerID.str());
    loader.loadData(path);
    ...
}

```

Figure 12: Code segment 1

Third, we write the code for query processing. As shown in Figure 13, we create an instance of *QueryProcessor* in *struct PPRGlobalIterateKernel*. As shown in Figure 13, in *struct PPRGlobalIterateKernel*, we define a variable with type *QueryProcessor*, create an instance of *QueryProcessor*, and pass in the current worker ID. The function *g\_func* and *dynamic\_inbound* are responsible for processing query. For these two functions, we call corresponding methods to access the shared memory and fetch the requested data.

```

struct PPRGlobalIterateKernel : public IterateKernel<int, double, PPRGlobal_data > {
    ...
    QueryProcessor* processor;
    ...
    PPRGlobalIterateKernel() {
        ...
        processor = new QueryProcessor(workerID.str());
        ...
    }
}

void g_func(const int& key, double& value, double& delta, PPRGlobal_data & data, vector<pair<int, double> *>* output) {
    ...
    int neibor_size = processor->get_vecsize(key_str.str().c_str(), "3");
    ...
}

double dynamic_inbound(const int& key, double dv, double maxdv, double sumdv, double hopbound, PPRGlobal_data & data) {
    ...
    float globalScore = processor->getFloat(key_str.str().c_str(), "1");
    float incomingDegree = processor->getFloat(key_str.str().c_str(), "2");
    ...
}

```

Figure 13: Code segment 2

Lastly, we add an *if* statement in method *read\_graph* to determine whether the present job is a data loading job or a query processing job. As shown in Figure 14, if the current job is responsible for data loading, it will perform the tasks of data loader. If the current job is for query processing, it will perform the tasks of query processor.

```

void read_graph(string& partition_file, TypedGlobalTable<int, double, double, PPRGlobal_data *> *table) {
    //...
    if(FLAGS_is_dataLoader) {
        DataLoader loader(workerID.str());
        loader.loadData(path);
    } else {
        //Process the query.
    }
    //...
}

```

Figure 14: Code segment 3

With the application of ParQ, we can start a data loader job to enable static data sharing and  $n$  query processors to process  $n$  queries concurrently.

## 4.2 Application on the top-k star query system

The ParQ system is also applied to the top-k PPR system. In the following sections, we will first introduce the top-k PPR system and then present how to apply ParQ to it.

### 4.2.1 System specification

The top- $k$  star query system answers pattern match queries in massive information networks. A massive information network describes labeled entities and the relationships among the entities. An entity can be a person or a thing with one attribute, type. On a massive information network, a star query aims to identify an unknown entity based on known facts. When representing the network as a graph, the star query problem can be modeled as a pattern matching problem.

In this problem, a query describes a known node, the type of an unknown node, and their relationship. The top- $k$  star query system aims to find the best  $k$  answers for a query from a graph. The candidates are ranked by their relevance to the known nodes. The relevance score between two nodes is computed by

$$\varphi(u, v) = \begin{cases} 1 & u = v \\ \min\{N, n_{u,v}\} \cdot \alpha^{l_{u,v}} & \text{otherwise} \end{cases},$$

where  $\alpha$  is a damping factor having the value between 0 and 1.  $l_{u,v}$  is the length of the shortest path between node  $u$  and  $v$ .  $n_{u,v}$  is the number of the shortest paths.  $N$  is a constant used to bound the value of  $n_{u,v}$ . That is, the relevance between two nodes is quantified by the length of the shortest path and the number of the shortest paths between them. The top- $k$  star query system uses breadth-first search to traverse a target graph to compute  $n_{u,v}$  and  $l_{u,v}$ . It also adopts a bounding technique to speed up the computation. The bounding technique can detect the top- $k$  answers without having to compute converged scores.

To answer a query, the top- $k$  star query system will start a process, or job, that loads the graph into memory and performs multiple BFSs on the target graph to measure relevance scores of candidates until the top- $k$  answers are obtained.

#### 4.2.2 Implementation process

As described in the last section, the top- $k$  star query system starts a job to answer a query. The job loads the target graph into memory and performs multiple BFSs on the in-memory target graph to measure relevance scores of candidates until the top- $k$  answers are obtained.

When answering different queries on the same target graph, the graph is the same, or static, and can be shared among different jobs. Therefore, we apply the ParQ system to the top- $k$  star query system to enable in-memory static data sharing.

First, we represent the target graph in a text file to apply the ParQ system, as shown in Figure 15. Each node is represented in one line. The first column is an integer representing the node name or the node key. The second column is the node type. The third column lists the node's neighbors. As explained in Section 3.2, the first column is always an integer and used as index. The labels and the data types of the following columns are specified in the first two lines.

```

type    neighbors
int     vector<int>
3177  -100  3098 2407 3052 5437 2073 3310 4382
3173  -100  3172
3169  -100  2441 3973 4194 6805 3702 3874 3900 5962 1746 4454 1445 6860 3168 2212
3097  -100  3096 3205 3009 4669 6430

```

Figure 15: A part of an item graph for top-k star query system

Second, we modify the code for data loader. Method *initDataGraph* in class *BoundWorker* is responsible for graph initializing. As shown in Figure 16, in this function we create an instance of *DataLoader* and pass in the current worker ID. Then, we call method *loadData* to create a shared memory and load the static data into the memory.

```

void BoundWorker::initDataGraph(){
    /**
     * @brief Create a shared memory and load the static data into the memory.
     */
    DataLoader loader(workerID.str());
    loader.loadData(path.str());
}

```

Figure 16: Code segment 1

Third, we write the code of query processor. Method *runIter* in class *BoundWorker* is responsible for computations. As shown in Figure 17, we define a variable with type *QueryProcessor* in the header file DBRWorker.h. We create an instance of *QueryProcessor*, pass in the current worker ID in *setMaiter*, and call functions in *runIter* to connect to the shared memory and get requested data.

```

class BoundWorker : public DSMKernel, public NetworkController{
    /**
     * @brief Create a shared memory and load the static data into the memory.
     */
    QueryProcessor* processor;
};

void BoundWorker::setMaiter(MaiterKernel<uint, BoundList, TargetNode*>* inmaiter) {
    /**
     * @brief Create a shared memory and load the static data into the memory.
     */
    processor = new QueryProcessor(workerID.str());
}

```

```

void BoundWorker::runIter(int tabId, const uint& k) {
    /**
     * ...
     * QueryProcessor* p = this->processor;
     * int node_type = p->get_int(key.str().c_str(), "1");
     * int neibor_size = p->get_vecsize(key.str().c_str(), "2");
     * ...
    }
}

```

**Figure 17: Code segment 2**

Lastly, as shown in Figure 18, we add an if statement in the function *Run* to determine whether the current job is a data loader or a query processor. If the current job is a data loader, it will load the static data and terminate. If the current job is a query processor, it will skip data loading and start to process query.

```

void DBRMaster::Run(Master* master){
    /**
     * ...
     * if(FLAGS_is_dataLoader){
     *     //Load the static data.
     * }else{
     *     //Process the query.
     * }
     * ...
    }
}

```

**Figure 18: Code segment 3**

With the application of ParQ, we can start a data loading job and  $n$  query processing job to process  $n$  queries concurrently.

### 4.3 Application on the belief propagation system

In order to further study the performance of the ParQ system, we also applied it to the Belief Propagation system. The following sections introduce the introduction of that system and present how to apply ParQ to it.

#### 4.3.1 System specification

The belief propagation system performs approximate inference on probabilistic graphical models. Probabilistic graphical models use a graph-based representation to capture uncertainty in real-world applications. In this graphical representation, the nodes correspond to the variables in the

real world. The edges represent direct probabilistic interactions between them. The query describes a set of variables

In this example, we consider a typical probabilistic graphical models, factor graphs, since any other graphical models can be converted to factor graphs. There are two types of nodes in factor graphs: variable nodes and factor nodes. Each variable node represents a random variable and each factor node describes a specific function that maps variable nodes to non-negative real-valued number.

The belief propagation system here adopts the sum-product algorithm, which answers query by computing marginal probabilities of factor graphs. In order to reduce scheduling cost, the belief propagation system calculates the priority of nodes and selects a set of messages to update at a time via the priority. In each steps, from the chosen node, the computation reaches the out-neighbors. It keeps propagating messages in both directions along edges until a stable situation is reached.

The belief propagation system extends the maiter framework [5]. In this system, the message is computed by accumulative computation in an asynchronous manner. A fraction of nodes are selected to update their scores in each pass. The chosen nodes will update its value ( $v$ ) and propagate the change of value ( $\Delta v$ ) to their neighbors.

To process a query, the belief propagation system will start a process, or job, that loads the factor graph into memory and performs an accumulative computation on the in-memory graph to measure the value of variables. The computation is performed until it reaches a stable situation.

#### 4.3.2 Implementation process

As descripted in the last section, the belief propagation system starts a job to perform approximate inference on probabilistic graphical models. The job loads the target graph into memory and performs an accumulative computation on the in-memory graph to measure the value of variables until the values reach a stable situation.

When answering different queries on the same target graph, the graph is the same, or static, and can be shared among different jobs. Therefore, we apply the ParQ system to the top-k star query system to enable in-memory static data sharing.

First, we represent the target graph in a text file to apply the ParQ system, as shown in Figure 19. Each node is represented in one line. The first column is an integer representing the node name or the node key. The second column lists the node's neighbors. As explained in Section 3.2, the first column is always an integer and used as index. The labels and the data types of the following columns are specified in the first two lines.

Second, we modify the code for data loader. Method *read\_data* is responsible for graph initializing. As shown in Figure 20, in this function we create an instance of *DataLoader* and pass in the current worker ID. Then, we call method *loadData* to create a shared memory and load the static data into the memory.

```
void read_data() {
    ...
    DataLoader loader(workerID.str());
    loader.loadData(path.str());
    ...
}
```

**Figure 19: Code segment 1**

Third, we write the code of query processor. Method *c\_fun* and *u\_fun* are responsible for computations. As shown in Figure 17, we call functions in *runIter* to connect to the shared memory and get requested data.

```
QueryProcessor* p = this->processor;
int node_type = p->get_int(key.str().c_str(), "1");
int localstatenum = p->get_int(key.str().c_str(), "2");
int outstatenum = p->get_int(key.str().c_str(), "3");
int neibor_size = p->get_vecsize(key.str().c_str(), "4");
```

**Figure 20: Code segment 2**

Lastly, we add an if statement in the function *Run* to determine whether the current job is a data loader or a query processor. If the current job is a data loader, it will load the static data and terminate. If the current job is a query processor, it will skip data loading and start to process query.

With the application of ParQ, we can start a data loading job and  $n$  query processing job to process  $n$  queries concurrently.

## CHAPTER 5

### APPLICATION AND EVALUATION

#### 5.1 Overview

The ParQ system has been applied to the top- $k$  Personalized PageRank system and the top- $k$  star query system. In order to examine the performance, we designed and performed a series of experiments, which includes running time and memory usage. For one query system that uses ParQ, multiple queries can be processed in parallel with one target graph in memory. We compare ParQ with the original system running in two ways. First, the original system can process multiple queries in serial with one target graph in memory (or “ONE CPY” for short). Alternatively, it can process queries in parallel but with multiple graphs in memory (or “DUPL CPY” for short).

#### 5.2 Performance of the top- $k$ Personalized PageRank system using ParQ

We first evaluate the performance of the top- $k$  PPR system adopting ParQ. In the following sections, after a brief introduction of the top- $k$  PPR system, an experiment overview and experimental results are presented.

##### 5.2.1 Experiment settings

We performed the experiments on a local computer cluster. The local cluster consisted of four machines connected by a 1Gb Ethernet switch. Each machine had 16GB RAM and one 1.86GHz CPU.

We used a set of queries on a 916K-line graph. The graph is a web graph that describes the directed links between webpages. The queries are personal preferences.

### 5.2.2 Running time

Figure 21 shows the comparison of the running time of the approaches for completing queries. First, it can be seen that for ParQ and Parallel/multi-graph, as the number of queries increases, the running time increases slightly. This is because that there is enough free memory as shown in Figure 22, since the top-k PPR system requires less memory to process one query.

Second, the results show that running time of ParQ is smaller than ONE CPY. The first reason is that ParQ processes the queries concurrently while ONE CPY processes queries in serial. The second reason is ONE CPY uses a long time to load the static data before processing queries, while ParQ can directly use the static data in the shared memory.

Third, the results also show that the running time of ParQ is smaller than DUPL CPY. This is because of limited free memory. In order to processes  $n$  queries concurrently, DUPL CPY keeps  $n$  copies of the static data in memory. When there is limited free memory, the running time of DUPL CPY increases due to increased workload. The second reason is that DUPL CPY also uses a long time to load the static data for each query.

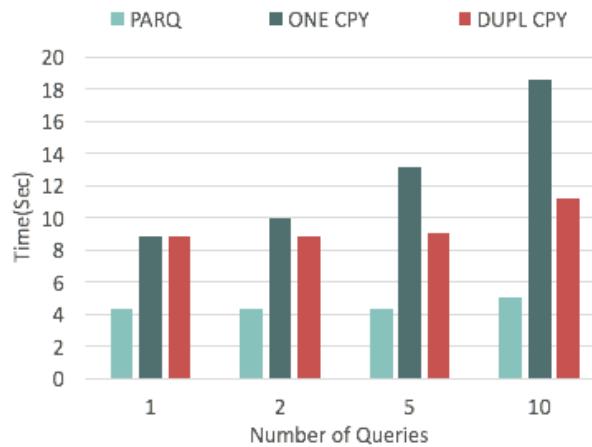


Figure 21: Running time for completing all queries

### 5.2.3 Memory usage

Figure 22 shows the comparison of the memory usage of the three approaches. First, the results show that ParQ requires less memory than DUPL CPY because ParQ keeps only one copy of the static data in memory while DUPL CPY needs  $n$  copies to process  $n$  queries.

Second, it can be seen that ParQ requires more memory than ONE CPY. Since each query job needs to maintain a set of dynamic data generated during computations, ParQ keeps  $n$  sets of dynamic data in memory to process  $n$  queries concurrently, while ONE CPY keeps only one set of dynamic data at a time because it processes  $n$  queries in serial.

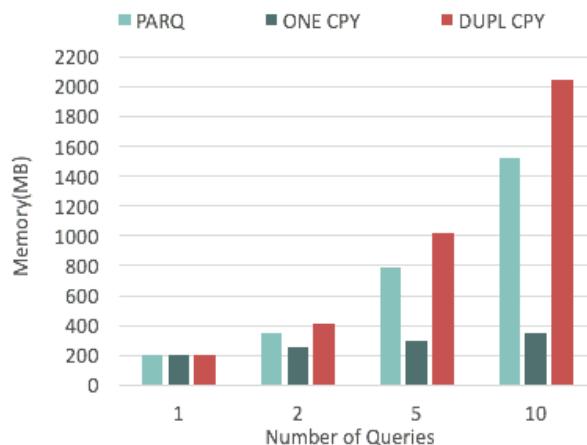


Figure 22: Memory usage

### 5.3 Performance of the top- $k$ star query system using ParQ

The proposed approach is also applied and evaluated on the top- $k$  star query system. What shall be introduced in later sections are a brief system introduction, an experiment overview, and experiment results.

### 5.3.1 Experiment settings

We performed the experiments on the same local computer cluster. The local cluster consisted of four machines connected by a 1Gb Ethernet switch. Each machine had 16GB RAM and one 1.86GHz CPU.

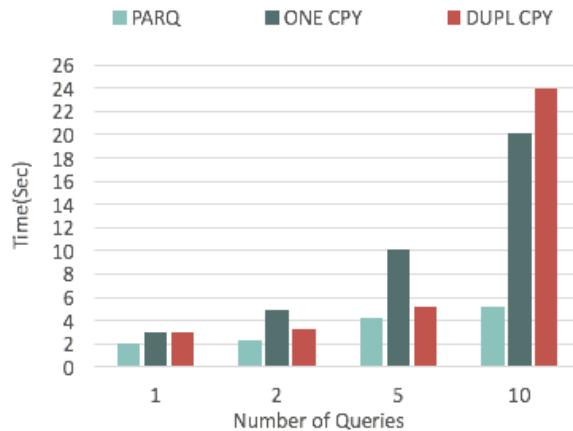
We used a set of queries on a 548K-line graph. The graph describes the collaboration among authors. The queries consisted of a known node set, an unknown node, and their relationship.

### 5.3.2 Running time

Figure 23 shows the performance comparison of the three approaches, focusing on the running time for completing various numbers of queries. First, the results show that for ParQ and Parallel/multi-graph, as the number of queries increases, the running time increases as well due to increased workload.

Second, it can be seen that the running time of ParQ is smaller than ONE CPY. This is because ParQ processes the queries concurrently while ONE CPY processes queries one by one.

Third, the results show the running time of ParQ is also smaller than DUPL CPY. When there are 10 queries, the running time of DUPL CPY is large. The reason is that although DUPL CPY processes 10 queries concurrently, it keeps 10 copies of the static data in memory. When there is limited free memory, the running time of DUPL CPY increases sharply.

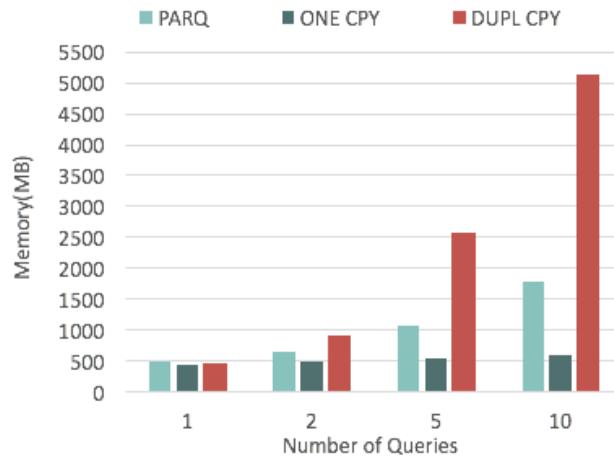


**Figure 23: Running time for completing all queries**

### 5.3.3 Memory usage

Figure 24 shows the performance comparison of the approaches, focusing on the memory usage for completing different numbers of queries. First, the results show that ParQ uses less memory than DUPL CPY, because ParQ keeps only one copy of the static data in memory while DUPL CPY requires  $n$  copies to process  $n$  queries.

On the other hand, the results show that ParQ requires more memory than ONE CPY. Each query job needs to maintain a set of dynamic data generated during computations. Since ParQ processes  $n$  queries concurrently, it keeps  $n$  sets of dynamic data in memory, while ONE CPY keeps only one set of dynamic data at a time because it processes  $n$  queries in serial.



**Figure 24: Memory usage**

## 5.4 Performance of the belief propagation system

At last, the proposed approach is applied and evaluated on the belief propagation system. What shall be introduced in later sections are a brief system introduction, an experiment overview, and experiment results.

### 5.4.1 Experiment settings

We performed the experiments on the same local computer cluster. The local cluster consisted of four machines connected by a 1Gb Ethernet switch. Each machine had 16GB RAM and one 1.86GHz CPU.

We used a set of queries on a 908K-line graph. The graph describes direct probabilistic interactions between given real facts.

### 5.4.2 Running time

Figure 25 shows the job completion time comparison of the three approaches. First, the results show that for ParQ and DUPL CPY approach, as the number of queries increases, the running time increases as well due to increased workload.

Second, it can be seen that the running time of ParQ is less than which of ONE CPY. The reason is that ParQ processes the queries concurrently while ONE CPY processes queries one by one.

Third, the results also show that the running time of ParQ is less than that DUPL CPY as well. Especially when there are 10 queries, the running time of DUPL CPY is large. The reason is that although DUPL CPY processes 10 queries concurrently, it keeps 10 copies of the static data in memory. When there is limited free memory, the running time of DUPL CPY increases sharply.

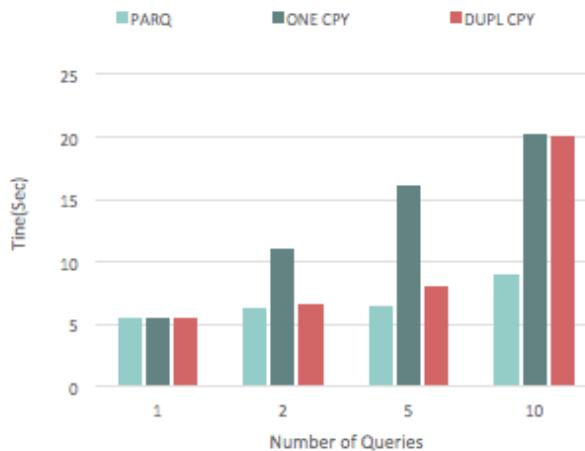


Figure 25: Running time for completing all queries

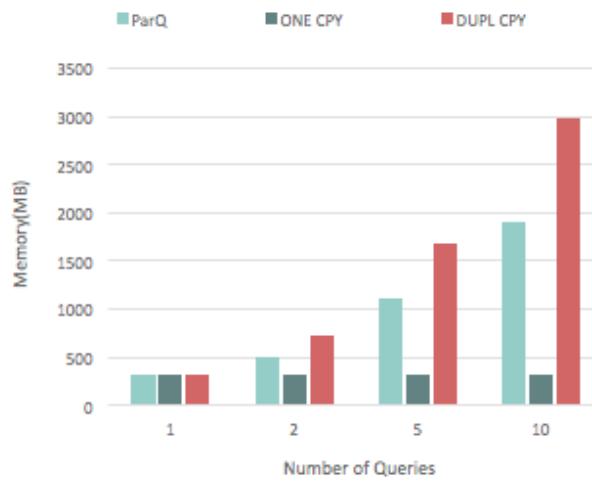
#### 5.4.3 Memory usage

Figure 26 shows the performance comparison of the approaches, focusing on the memory usage for completing different numbers of queries. Similarly, first, the results show that ParQ uses less memory than DUPL CPY, because ParQ keeps only one copy of the static data in memory while DUPL CPY requires  $n$  copies to process  $n$  queries.

On the other hand, the results show that ParQ requires more memory than ONE CPY. Each query job needs to maintain a set of dynamic data generated during computations. Since ParQ processes  $n$  queries concurrently, it keeps  $n$  sets of dynamic data in memory, while ONE CPY keeps only one set of dynamic data at a time because it processes  $n$  queries in serial.

However, compared with the performance of the PageRank system and the star query system, ParQ here requires relative much more memory than ONE CPY and is slightly better than DUPL CPY. Recall that the belief propagation system calculates the priority of nodes and selects a set of messages to update at a time via the priority. In each steps, from the chosen node, the computation reaches the out-neighbors. It will visit all graph nodes before the computation reaches a stable situation is reached while the previous two systems will visit a small fraction of nodes. The dynamic data of each query processing job will be relative much more.

As a conclusion, ParQ is more suitable in the situations when the size of static data is much larger than the size of dynamic data, although it is still useful in the opposite situations.



**Figure 26: Memory usage**

## CHAPTER 6

### CONCLUSION AND FUTURE WORK

The following table provides an overview of the plan for completion of the research including what I have already done.

**Table 1: Timeline**

| Activity   | Period                |
|--|-----------------------|
| Read papers and do research.                             | Sep. 2014 – Oct. 2014 |
| Formulate overall plan.                                  | Nov. 2014             |
| Evaluate possible approaches that may achieve the goals. | Dec. 2014             |
| Formulate the overall design of ParQ.                    | Jan. 2015             |
| Implement main methods of ParQ.                          | Feb. 2015 – Apr. 2015 |
| Improve ParQ to adapt to distributed environment.        | May. 2015 – Jun. 2015 |
| Complete writing and proposal.                           | Jul. 2015 – Sep. 2015 |
| Study more query systems.                                | Oct. 2015 – Dec. 2015 |
| Improve ParQ so that it can process more data types.     | Jan. 2016 – Feb. 2016 |
| Adopted improved ParQ and perform experiments            | Mar. 2016 – Apr. 2016 |
| Complete writing   | May. 2016             |
| Complete thesis defense.                                 | Jul. 2016             |

## CHAPTER 7

### RELATED WORK

#### 7.1 Large structured dataset processing

Large structured dataset processing has been studied in various application areas. Many projects have been developed for processing structured data, such as Pregel [7] and Giraph [8]. Based on BSP model [9], they implement a vertex-centric computation to process large graphs on parallel machines or multi-cores. However, they focus on only graph data and suffer the problem of memory inefficient when running concurrent jobs. There are other systems try to improve the computation. For example, Seraph [3] is a graph processing system that enables efficient job-level parallelism. However, it is an implementation framework and cannot be directly applied to existing systems. Our system allows multiple concurrent jobs to share data in memory and can be easily used by existing systems.

#### 7.2 Algorithm implementations

Recently, many algorithms have been proposed to discover information from structured data such as a fast approach top- $k$  path-based relevance query [4] and an index-free approach for top- $k$  star queries [5]. Many implementations of algorithms are designed for single job execution, which are inefficient when running multiple concurrent jobs because of inefficient use of memory. Our system makes it possible to easily enable job-level parallelism for different implementations.

## References

- [1] Knowledge Graph. [https://en.wikipedia.org/wiki/Knowledge\\_Graph](https://en.wikipedia.org/wiki/Knowledge_Graph)
- [2] Freebase. <https://en.wikipedia.org/wiki/Freebase>
- [3] Xue, J., et al.: Seraph: an efficient, low-cost system for concurrent graph processing. Proceedings of the 23rd international symposium on High-performance parallel and distributed computing. ACM (2014)
- [4] Khemmarat, S., et al.: Fast top-k path-based relevance query on massive graphs. ICDE'14 (2014)
- [5]
- [6] Zhang, Y., et al.: Accelerate large-scale iterative computation through asynchronous accumulative updates. pp. 13–22. ScienceCloud '12, ACM (2012)
- [7] Malewicz, G., et al.: Pregel: A system for large-scale graph processing. In SIGMOD '10 (2010)
- [8] Giraph. <http://giraph.apache.org/>
- [9] Valiant, L.: A bridging model for parallel computation. Commun. ACM, 33(8):103–111 (Aug. 1990)